



(https://swayam.gov.in)



(https://swayam.gov.in/nc\_details/NPTEL)

vijay.mahawar@hotmail.com

NPTEL (https://swayam.gov.in/explorer?ncCode=NPTEL) » Introduction to Large Language Models (LLMs)  
(course)

Click to register  
for Certification  
exam

(https://examform.nptel.ac.in/2025\_01/exam\_form/dashboard)

If already  
registered, click  
to check your  
payment status

Course  
outline

About NPTEL  
( )

How does an  
NPTEL online  
course work?  
( )

Week 1 ( )

Week 2 ( )

Week 3 ( )

Week 4 ( )

Week 5 ( )

Week 6 ( )

Week 7 ( )

Week 8 ( )

Week 12 : Assignment 12

The due date for submitting this assignment has passed.

Due on 2025-04-16, 23:59 IST.

Assignment submitted on 2025-04-08, 19:19 IST

1) Which statements correctly characterize “bias” in the context of LLMs? 1 point

- 1. Bias can generate objectionable or stereotypical views in model outputs.
- 2. Bias is always intentionally introduced by malicious data curators.
- 3. Bias can cause harmful real-world impacts such as reinforcing discrimination.
- 4. Bias only affects low-resource languages; high-resource languages are unaffected.

- ☐ 1 and 2
- ☒ 1 and 3
- ☐ 2 and 4
- ☐ 1, 3, and 4

Yes, the answer is correct.  
Score: 1

Accepted Answers:  
1 and 3

2) The Stereotype Score (ss) refers to: 1 point

- ☐ The frequency with which a language model rejects biased associations.
- ☐ The measure of how often a model’s predictions are meaningless as opposed to meaningful.
- ☐ A ratio of positive sentiment to negative sentiment in model outputs.
- ☒ The proportion of examples in which a model chooses a stereotypical association over an anti-stereotypical one.

Yes, the answer is correct.  
Score: 1

Accepted Answers:  
The proportion of examples in which a model chooses a stereotypical association over an anti-stereotypical one.

## Week 9 ()

## Week 10 ()

## Week 11 ()

## Week 12 ()

● Lec 36 :  
Responsible  
LLMs (unit?  
unit=105&lesso  
n=106)

● Lec 37 :  
Conclusion:  
Expert Panel  
Discussion  
(unit?  
unit=105&lesso  
n=107)

● Lecture Material  
(unit?  
unit=105&lesso  
n=108)

○ Feedback Form  
(unit?  
unit=105&lesso  
n=109)

● Quiz: Week 12  
: Assignment  
12  
(assessment?  
name=110)

## Year 2025 Solutions ()

3) Which of the following are prominent sources of bias in LLMs?

1 point

1. Improper selection of training data leading to skewed distributions.
2. Reliance on older datasets causing “temporal bias.”
3. Overemphasis on low-resource languages causing “linguistic inversion.”
4. Unequal focus on high-resource languages resulting in “cultural bias.”

☐ 1 and 2 only

☐ 2 and 3 only

☒ 1, 2, and 4

☐ 1, 3, and 4

Yes, the answer is correct.

Score: 1

Accepted Answers:

1, 2, and 4

4) In the context of bias mitigation based on adversarial triggers, which best describes the goal of prepending specially chosen tokens to prompts? 1 point

☐ To directly fine-tune the model parameters to remove bias

☐ To override all prior knowledge in a model, effectively “resetting” it

☒ To exploit the model’s distributional patterns, thereby neutralizing or flipping biased associations in generated text

☐ To randomly shuffle the tokens so that the model becomes more robust

Yes, the answer is correct.

Score: 1

Accepted Answers:

To exploit the model’s distributional patterns, thereby neutralizing or flipping biased associations in generated text

5) Which of the following best describes the “regard” metric?

1 point

☐ It is a measure of how well a model can explain its internal decision process.

☐ It is a measurement of a model’s perplexity on demographically sensitive text.

☐ It is the proportion of times a model self-corrects discriminatory language.

☒ It is a classification label reflecting the attitude towards a demographic group in the generated text.

Yes, the answer is correct.

Score: 1

Accepted Answers:

It is a classification label reflecting the attitude towards a demographic group in the generated text.

6) Which of the following steps compose the approach for improving response safety via in-context learning? 1 point

☒ Retrieving safety demonstrations similar to the user query.

☐ Fine-tuning the model with additional labeled data after generation.

☒ Providing retrieved demonstrations as examples in the prompt to guide the model’s response generation.

☒ Sampling multiple outputs from LLMs and choosing the majority opinion.

No, the answer is incorrect.

Score: 0

Accepted Answers:

Retrieving safety demonstrations similar to the user query.

Providing retrieved demonstrations as examples in the prompt to guide the model’s response generation.

7) Which statement(s) is/are correct about how high-resource (HRL) vs. low-resource languages (LRL) affect model training? **1 point**

- ☐ LRLs typically have higher performance metrics due to smaller population sizes.
- ☒ HRLs get more data, so the model might overfit to HRL cultural perspectives.
- ☒ LRLs are often under-represented, leading to potential underestimation of their cultural nuances.
- ☒ The dominance of HRLs can cause a reinforcing cycle that perpetuates imbalance.

Yes, the answer is correct.

Score: 1

Accepted Answers:

*HRLs get more data, so the model might overfit to HRL cultural perspectives.*

*LRLs are often under-represented, leading to potential underestimation of their cultural nuances.*

*The dominance of HRLs can cause a reinforcing cycle that perpetuates imbalance.*

8) The “Responsible LLM” concept is stated to address: **1 point**

- ☐ Only the bias in LLMs
- ☒ A set of concerns including explainability, fairness, robustness, and security
- ☐ Balancing training costs with carbon footprint
- ☐ Implementation of purely rule-based safety filters

Yes, the answer is correct.

Score: 1

Accepted Answers:

*A set of concerns including explainability, fairness, robustness, and security*

9) Within the StereoSet framework, the icat metric specifically refers to: **1 point**

- ☐ The ratio of anti-stereotypical associations to neutral associations
- ☐ The percentage of times a model refuses to generate content deemed hateful
- ☐ A measure of domain coverage across different demographic groups
- ☒ A balanced metric capturing both a model’s language modelling ability and the tendency to avoid stereotypical bias

Yes, the answer is correct.

Score: 1

Accepted Answers:

*A balanced metric capturing both a model’s language modelling ability and the tendency to avoid stereotypical bias*

10) Bias due to improper selection of training data typically arises in LLMs when: **1 point**

- ☐ Data are selected exclusively from curated, balanced sources with equal representation
- ☐ The language model sees only real-time social media feeds without any historical texts
- ☒ The training corpus over-represents some topics or groups, creating a skewed distribution
- ☐ All data are automatically filtered to remove any demographic markers

Yes, the answer is correct.

Score: 1

Accepted Answers:

*The training corpus over-represents some topics or groups, creating a skewed distribution*