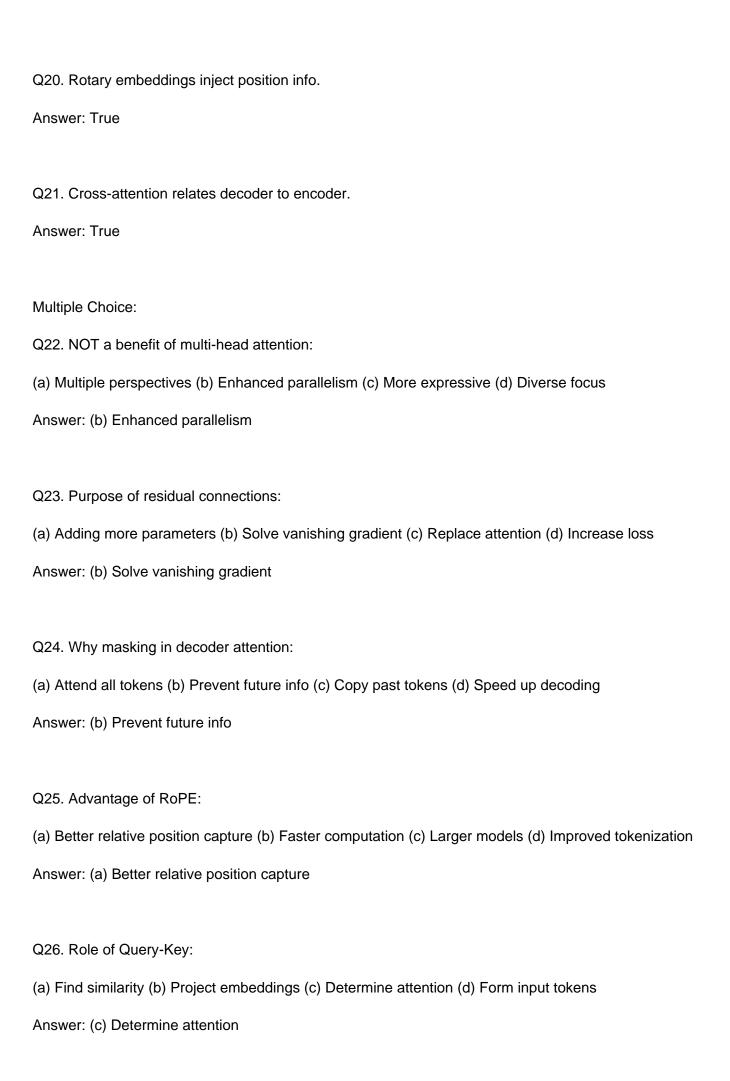
Combined Full Questions and Answers (Chapters 3 to 13)
Chapter 3-4: Statistical Language Models, Word2Vec, GloVe, TF-IDF, Bias
True/False:
Q1. Lemmatization is more computationally expensive than stemming.
Answer: True
Q2. Sigmoid function output is between -1 and 1.
Answer: False
Q3. SentencePiece requires tokenized input data.
Answer: True
Q4. Scalar multiplication is a non-linear operation.
Answer: False
Allower. False
Q5. Dependency parsing mainly focuses on word order.
Answer: False
Multiple Choice:
Q6. In dependency parsing, the head word is called:
(a) Prime (b) Root (c) Parent (d) Core
Answer: (a) Prime

Q7. Study of internal word structure is:
(a) Syntax (b) Phonology (c) Morphology (d) Semantics
Answer: (c) Morphology
Q8. In gradient descent, what is updated?
(a) Parameters (b) Loss (c) Features (d) Epochs
Answer: (a) Parameters
Q9. Dependency parsing reveals:
(a) Sentence semantics (b) Linear sentence form (c) Alphabetical order (d) Morphological structure
Answer: (d) Morphological
Chapter 5: Neural Language Models (RNN, LSTM, Attention)
True/False:
Q10. CNN filters act like n-gram detectors in NLP.
Answer: True
Q11. Padding helps uniform sentence length handling.
Answer: True
Q12. LSTMs use gates to manage flow of information.
Answer: True
Q13. Attention helps model focus on important input parts.

Q14. Fixed vector is the only output in attention models.
Answer: False
Multiple Choice:
Q15. Which combats vanishing gradients?
(a) CNN (b) RNN (c) LSTM (d) Vanilla NN
Answer: (c) LSTM
Q16. Hierarchical convolution helps in:
(a) Speed (b) Complexity handling (c) Activation (d) Optimizing pooling
Answer: (b) Complexity handling
Chapter 6: Transformers (Self-Attention, RoPE)
True/False:
Q17. Feed-forward networks introduce non-linearity.
Answer: True
Q18. Multi-query attention reduces computation.
Answer: True
Q19. Sliding window attention attends locally.
Answer: True

Answer: True



(Continuing in next batch due to long length)
Chapter 7: Language Model Pretraining (BERT, T5, GPT)
True/False:
Q27. Encoder-only models use Causal Language Modeling.
Answer: False
Q28. The [SEP] token is used to separate sentences.
Answer: True
Q29. Decoder-only models have separate encoders.
Answer: False
Q30. BERT uses masked tokens prediction.
Answer: True
Q31. T5 uses span-based corruption for pretraining.
Answer: True
Multiple Choice:
Q32. Special tokens used in BERT for sentence pairs:
(a) [PAD], [MASK] (b) [SEP], [CLS] (c) [UNK], [SEP] (d) [CLS], [PAD]
Answer: (b) [SEP], [CLS]

Q33. Pretraining objective for T5:
(a) Next Token Prediction (b) Translation (c) Span Corruption (d) MLM
Answer: (d) Span Corruption
Q34. BERT NSP task determines:
(a) Sentiment (b) Next Sentence Link (c) Word similarity (d) Masked tokens
Answer: (b) Next Sentence Link
Chapter 8: Fine-Tuning and Alignment (RLHF, DPO, CoT, ToT, GoT)
True/False:
Q35. CoT enables multiple reasoning paths.
Answer: True
Q36. ToT framework allows parallel path exploration.
Answer: True
Q37. Prompt engineering modifies model weights.
Answer: False
Q38. GoT forms clear sequential paths.
Answer: False
Q39. ToT uses DFS and BFS for problem solving.
Answer: True

Q40. Goal of prompt engineering:
a) Design effective prompts (b) Fine-tune LLMs (c) Build loss function (d) Tokenize input
Answer: (a) Design effective prompts
Q41. In-context learning implies:
(a) Model retraining (b) Learning from examples in prompts (c) RLHF application (d) Labeling
Answer: (b) Learning from examples in prompts
Q42. Example of prompting:
a) Summarize this text (b) Mask this token (c) Fine-tune model (d) Optimize decoder
Answer: (a) Summarize this text
Q43. How in-context learning helps:
a) Regularize model (b) Avoid retraining (c) Clean data (d) Pretrain
Answer: (b) Avoid retraining
Q44. CoT advantage:
a) Faster training (b) Multi-step reasoning (c) Shorter inputs (d) More loss
Answer: (b) Multi-step reasoning
Continuing next set for Chapters 9-13)
Chapter 9: Efficient Fine-Tuning (LoRA, Quantization, BitFit)

Multiple Choice:

True/False:
Q45. Pruning removes whole layers.
Answer: False
Q46. Mixed precision uses only FP32.
Answer: False
Q47. KL divergence used in distillation.
Answer: True
Q48. LoRA and BitFit help fine-tune large models.
Answer: True
Q49. Black-box knowledge distillation is possible.
Answer: True
Multiple Choice:
Q50. Hyperparameter balancing distillation and task loss:
(a) Beta (b) Alpha (c) Delta (d) Gamma
Answer: (b) Alpha
Q51. Soft labels refer to:
(a) Labels from fine-tuned model (b) Raw logits (c) Ground truth (d) Attention weights
Answer: (c) Ground truth
Q52. Benefit of quantisation:
(a) Faster training (b) Smaller model size (c) More parameters (d) Slower memory

Answer: (b) Smaller model size					
Q53. Which is NOT a prompt compression technique:					
(a) Context dropout (b) Prefix tuning (c) Keyword extraction (d) Prompt distillation					
Answer: (c) Keyword extraction					
Q54. Mixed precision achieves:					
(a) Balancing FP16 and FP32 (b) Only FP32 (c) Only FP16 (d) Slower inference					
Answer: (a) Balancing FP16 and FP32					
Chapter 10: Augmented Large Language Models (RAG)					
(True/False and MCQs will continue similarly)					